

RAG Chatbot Prototype Sample Document

Retrieval-augmented generation (RAG) combines a retriever with a language model so that answers stay grounded in real documents instead of relying purely on the model parametric memory.

This prototype stores document chunks and their embeddings in a Postgres table using the pgvector extension, and performs cosine similarity search to find the most relevant chunks for a query.

Azure OpenAI supplies both the text embeddings and the chat model used to generate the final answer.